

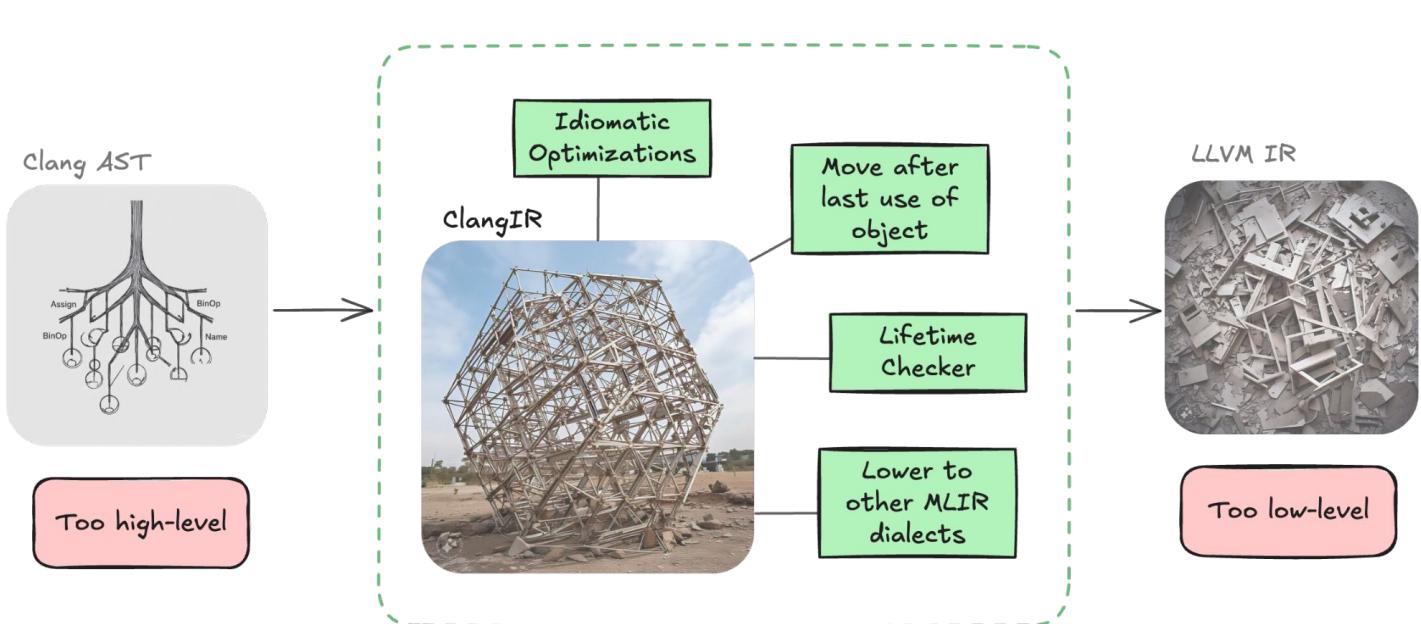
# ClangIR: Upstreaming an Incubator Project

Bruno Cardoso Lopes, Meta  
Andy Kaylor, NVIDIA



# ClangIR recap

Compiler pipeline



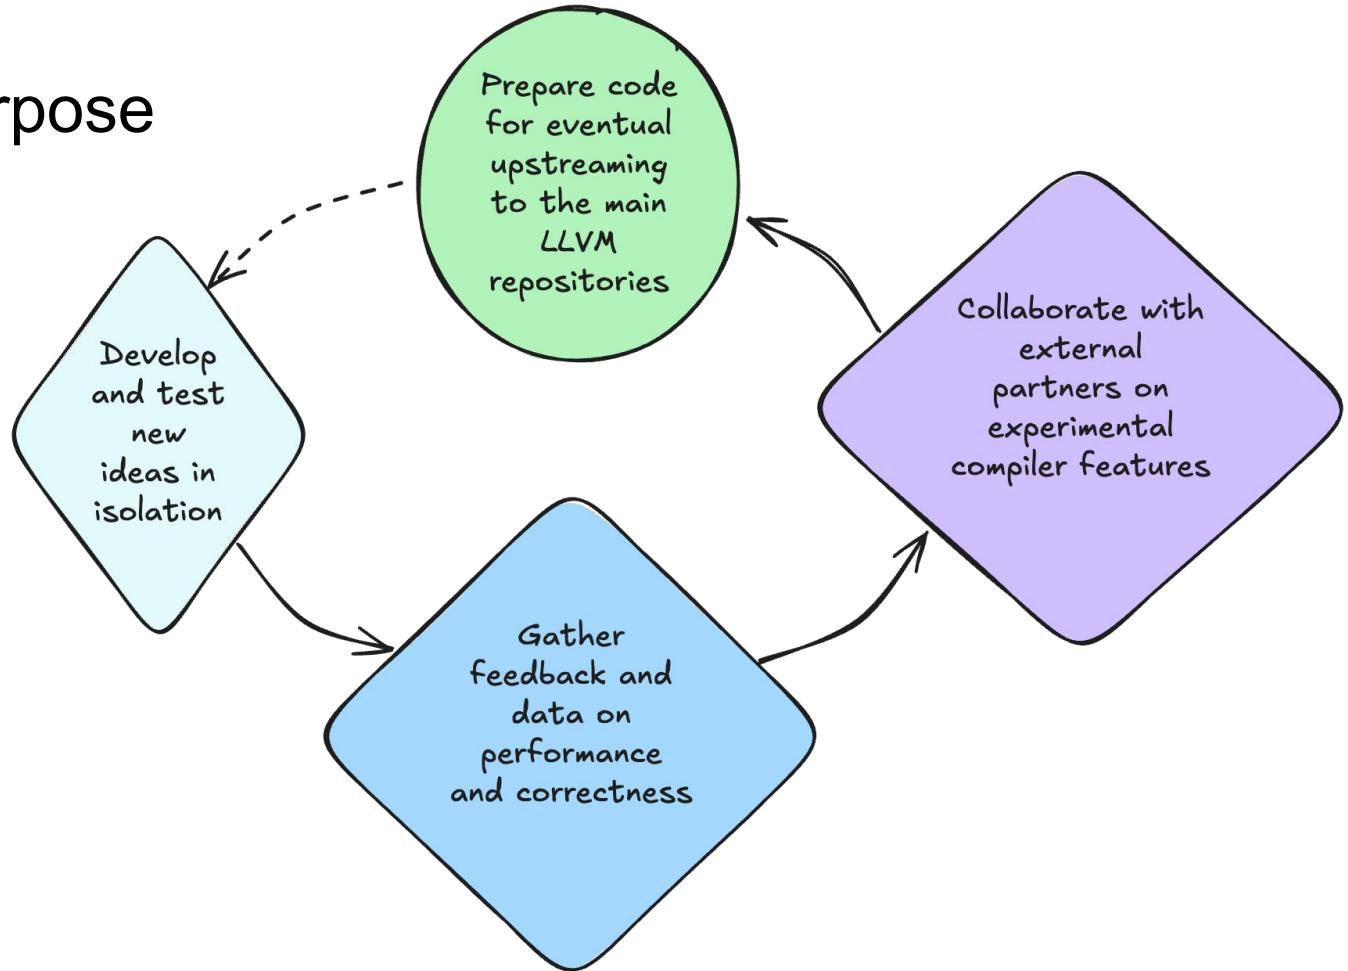
# ClangIR timeline

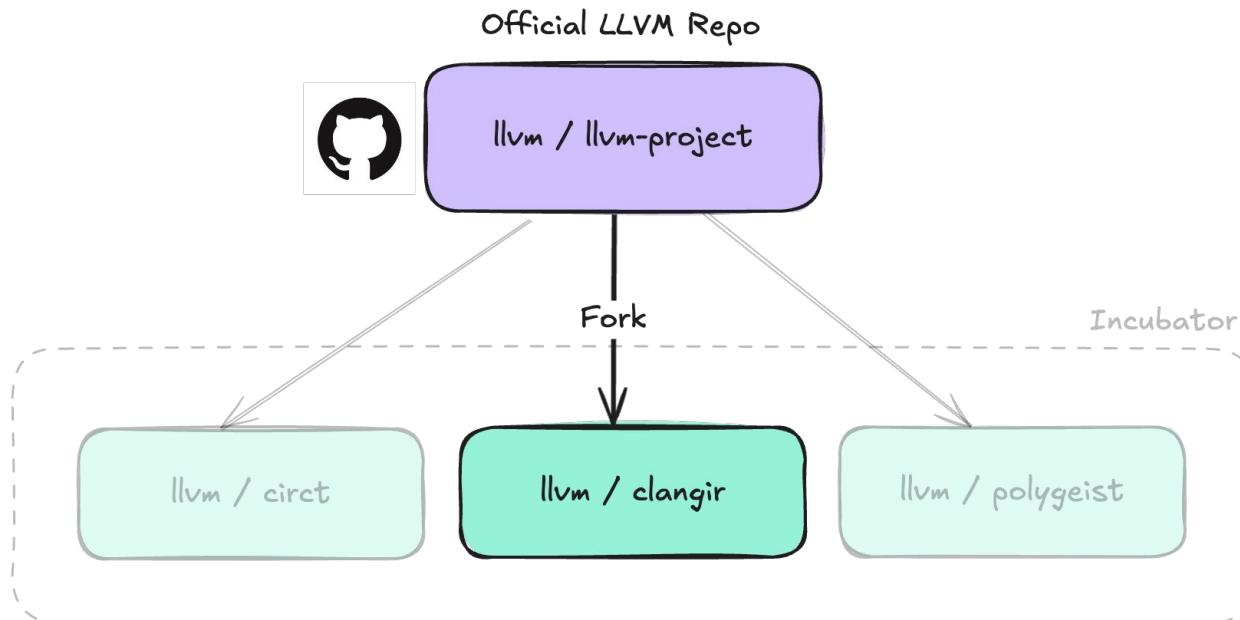


# ClangIR timeline



# Incubator's purpose





# From RFC to incubator

C/C++ MLIR frontend working group.

Move to LLVM incubator.



# Incubation advantages

Ability to use CI and other LLVM project infrastructure

Project visibility

Path for building upstreaming agreement

The LLVM umbrella attracts more contributors

# Incubation downsides

Less visibility than llvm-project

Users prefer llvm-project over a fork

Investment risk

Might never meet the bar

Never upstreamed

Incubation process is not formalized, consensus can be challenging

# Life in clangir incubator

[clangir.org](https://clangir.org) guides

LLVM practices enforced !

Paving the way for upstreaming

Additional freedom for exploration 





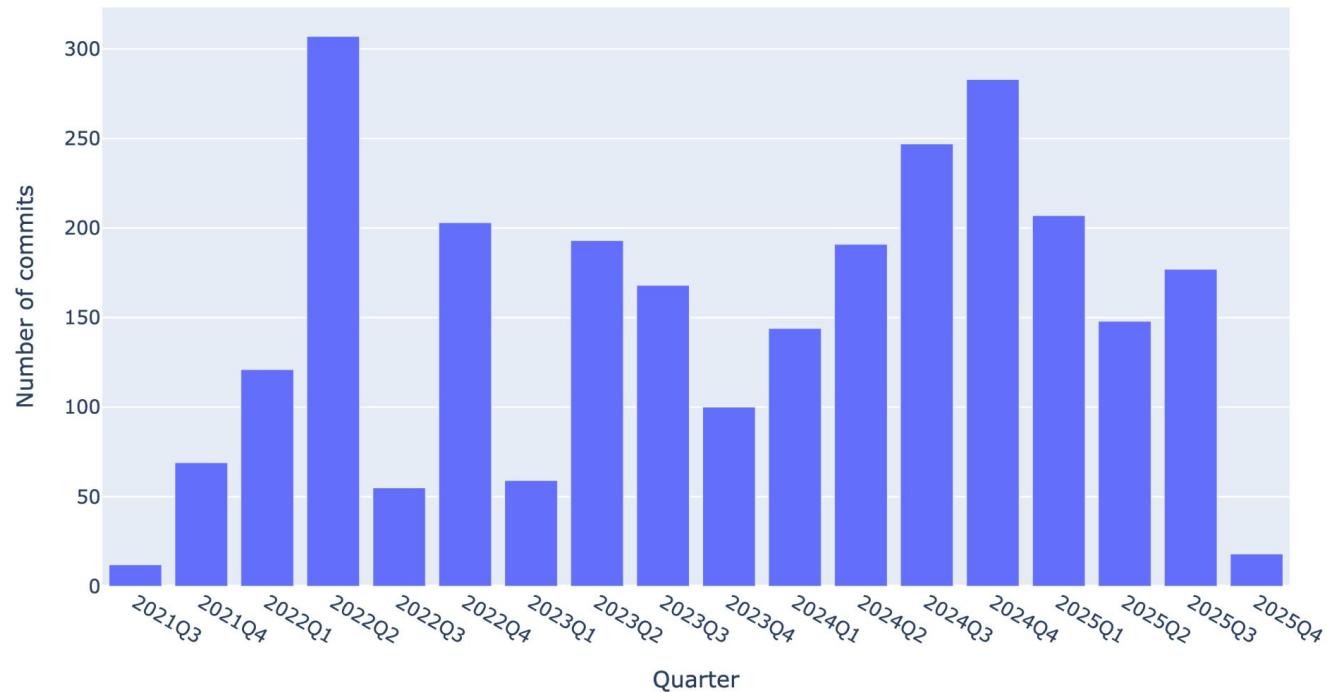
# Challenges

Wait! I heard this was working

*Always in motion is top-of-tree* 

Upstream and incubator differences

# Contributions



# Authors come and go ✈️

Contribution span varies

Good for incremental pieces (e.g. builtins)

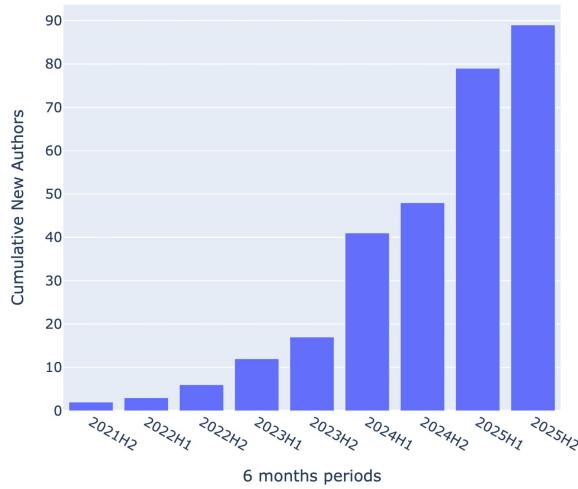
Especially attractive to newcomers

Bad for solidifying expertise

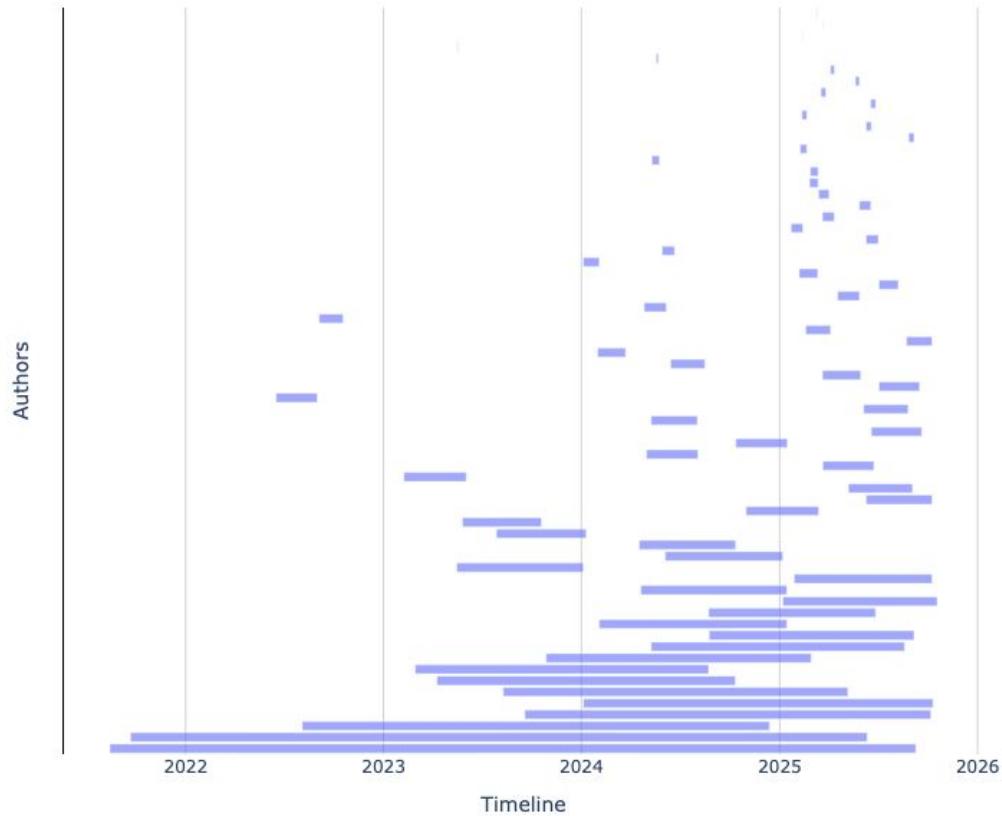
Some couldn't wait for upstreaming



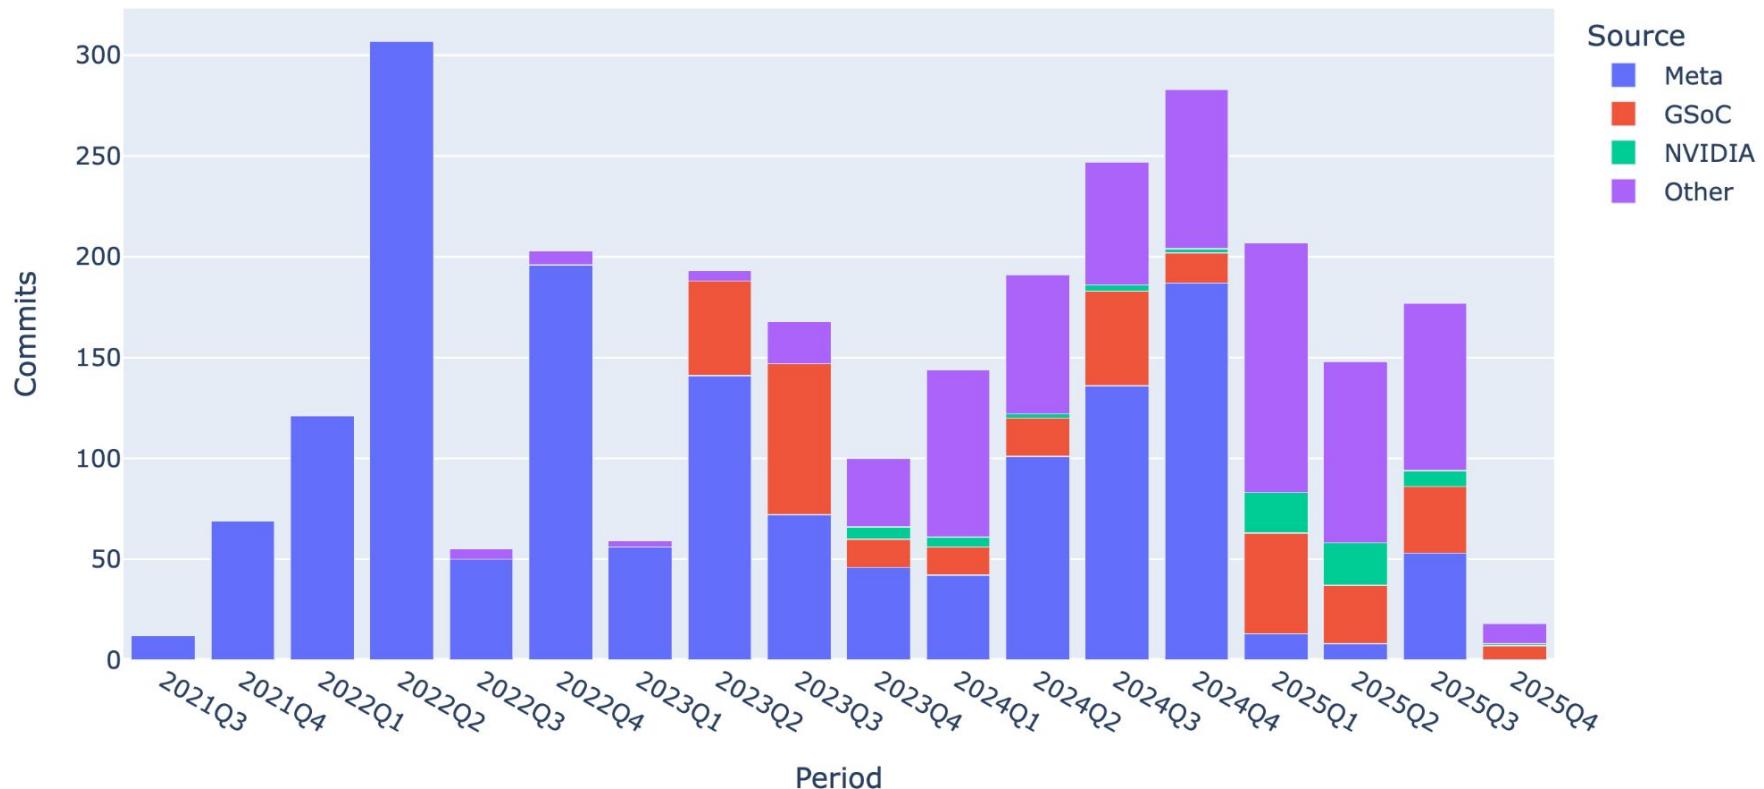
### Contribution Lifespans (per author)



23 only contributed once



## Incubator: Grouped Commits per Quarter



# Incubator - In hindsight

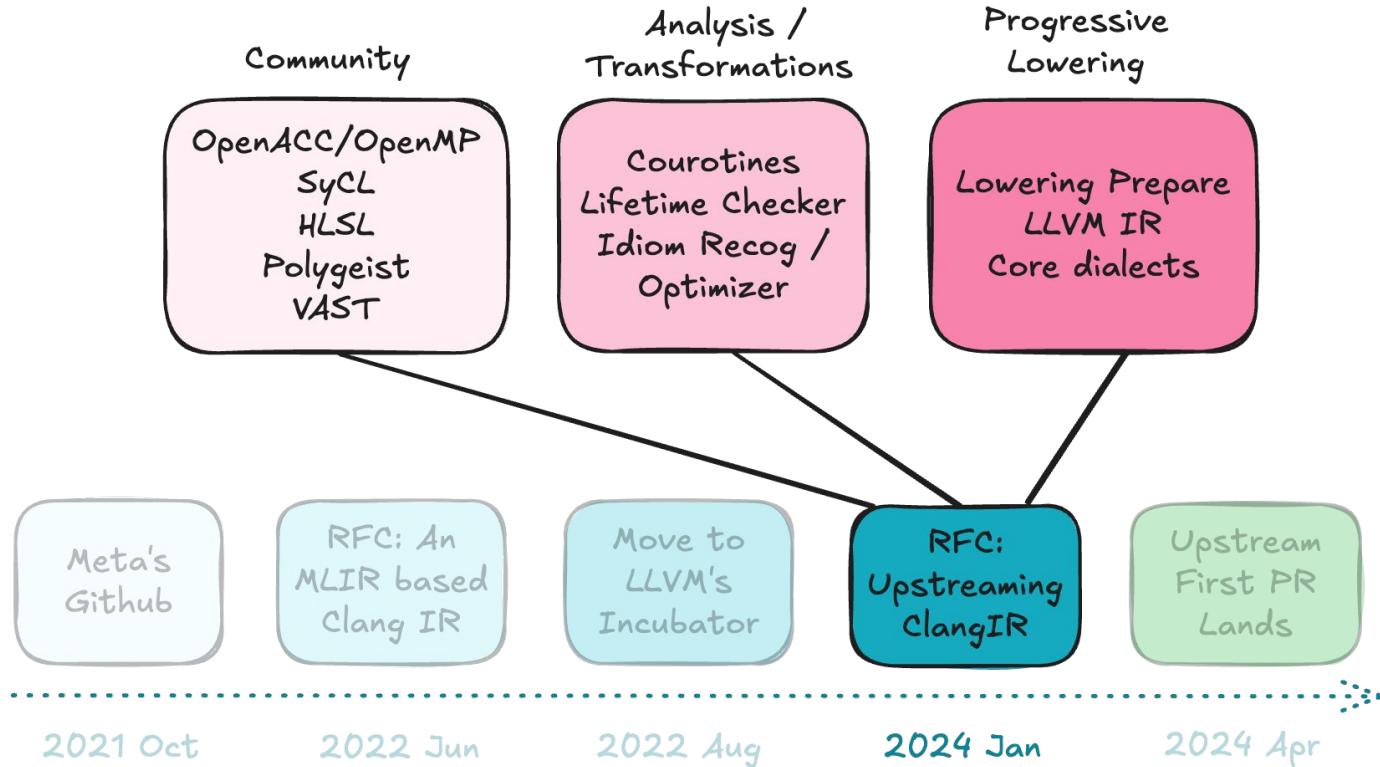
Challenging

... but worth it

Community up and running

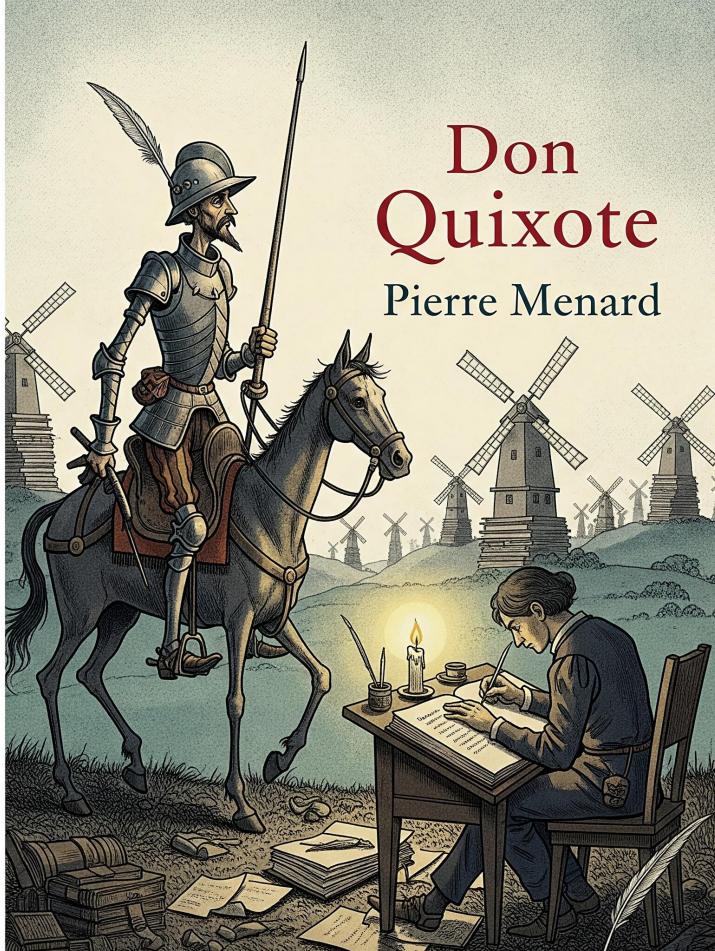
Steady progress in upstreaming





# Upstreaming begins



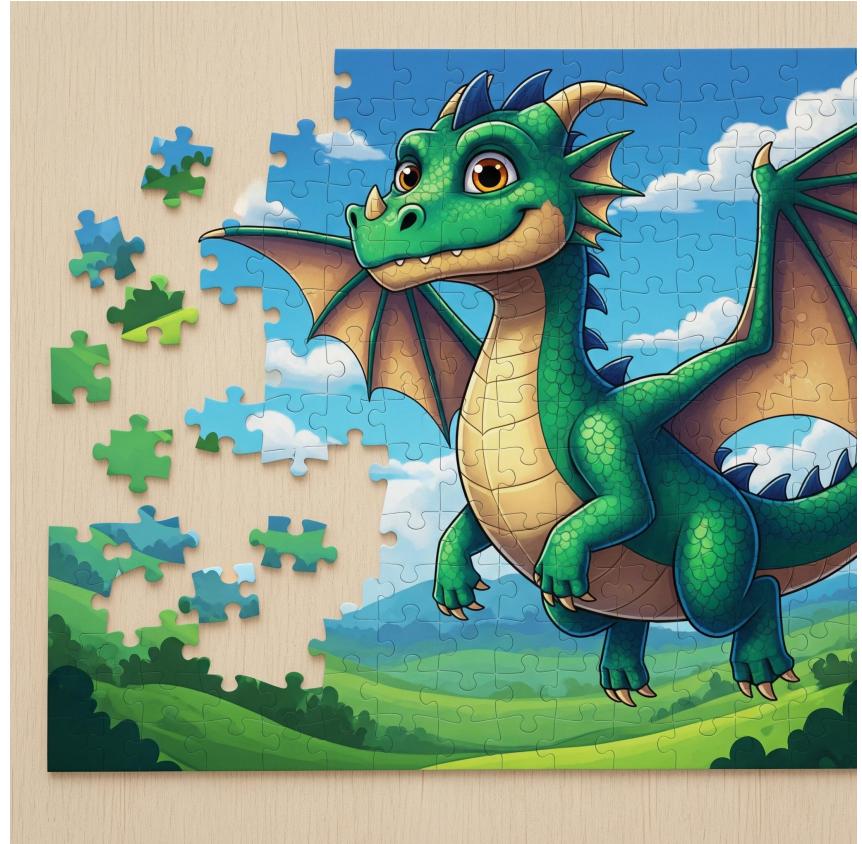


## A useful accident

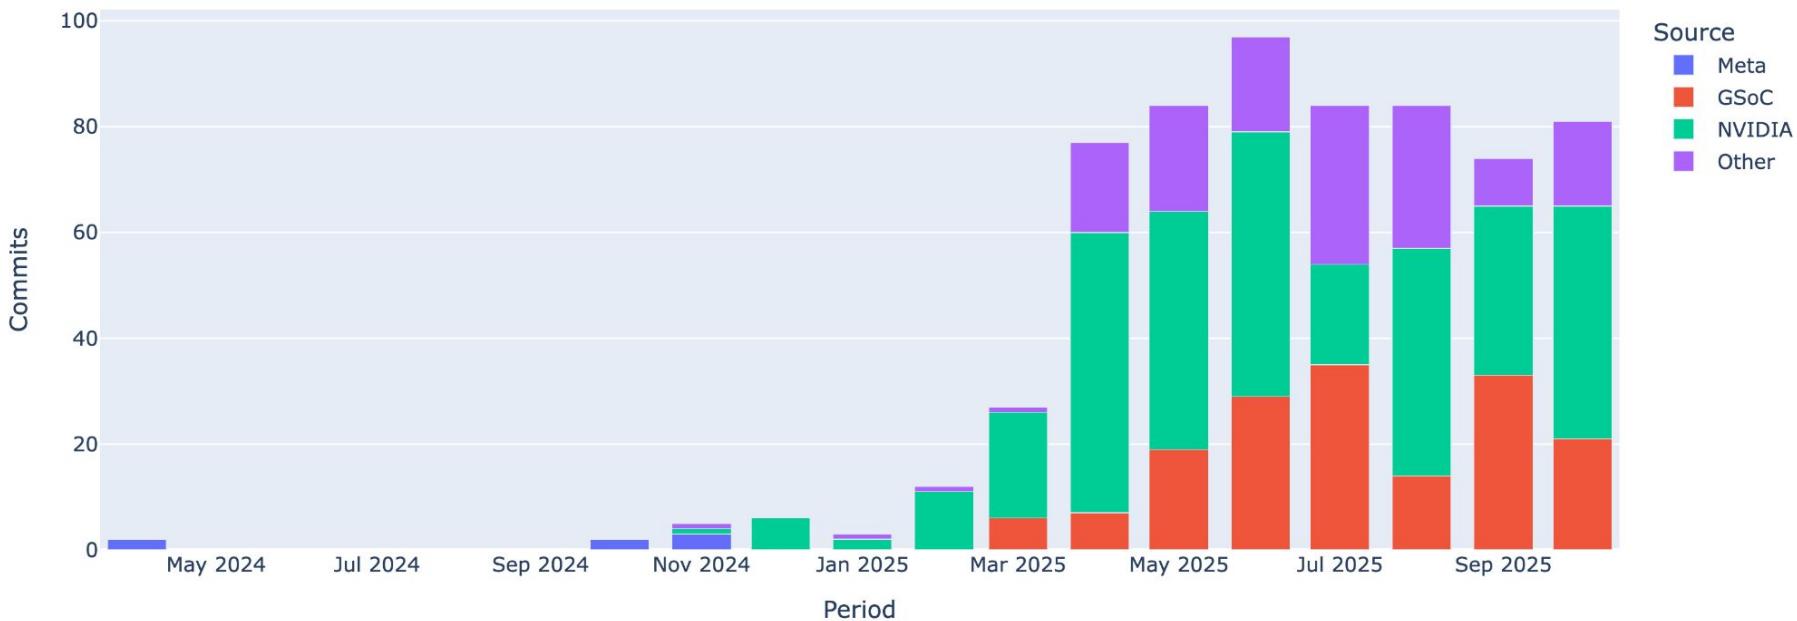
- The contributors to the upstreaming effort are mostly different than the original contributors to the incubator
- Pierre Menard: Author of the Quixote
  - New contributors creating the same code
  - Fresh perspectives bring new context
- Wax on, wax off
  - Learning by going through the motions
- Result: twice as many people have deep understanding of the code

# Going live!

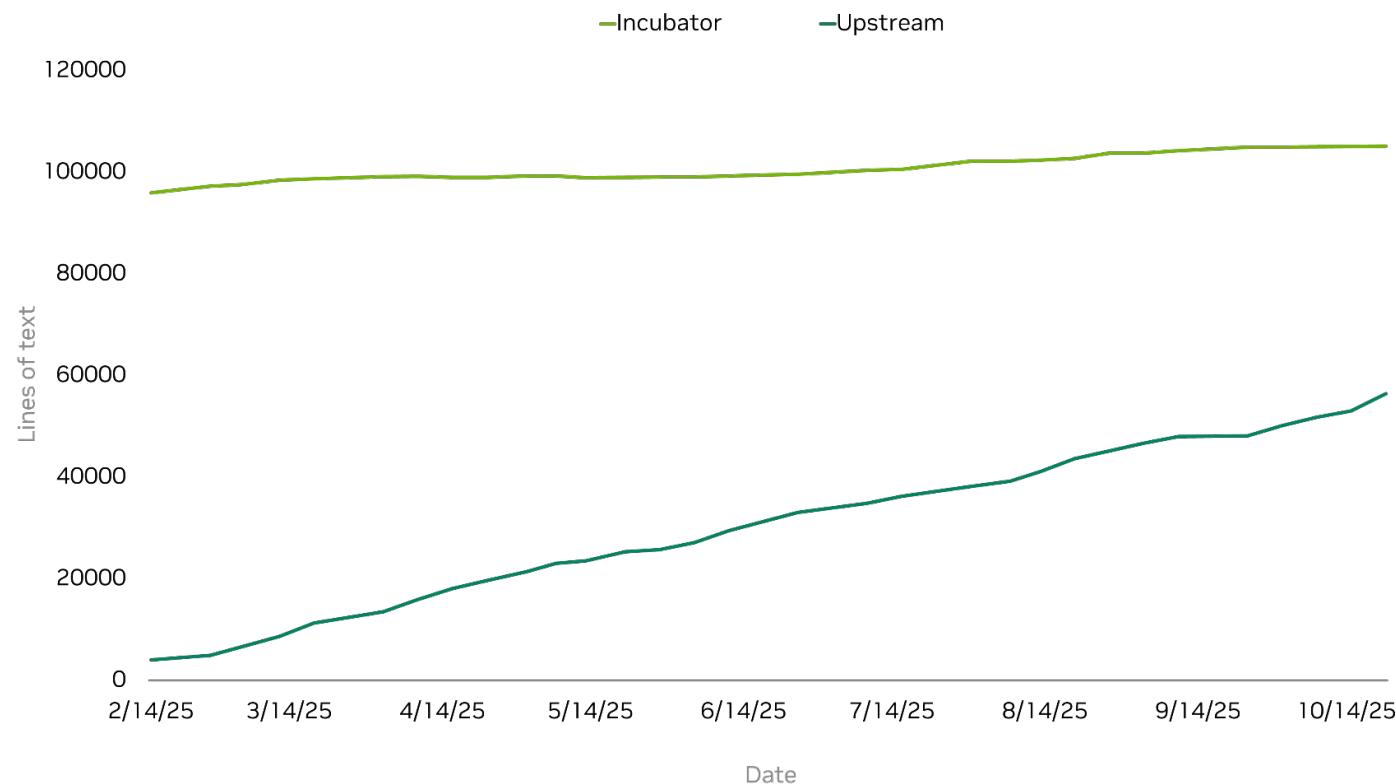
- Doing everything *as if* from scratch
  - The clang maintainers didn't want the incubator project merged en masse
  - There was no practical way to "replay" the incubator commits
- Each PR for upstreaming code is treated exactly as if it were a new contribution
- Each PR is held to the same standards as any LLVM contribution
  - Contributor must thoroughly understand the code
  - All functionality must be tested
  - Conformance to coding standards
  - Single purpose
  - Isolated change, as small as possible



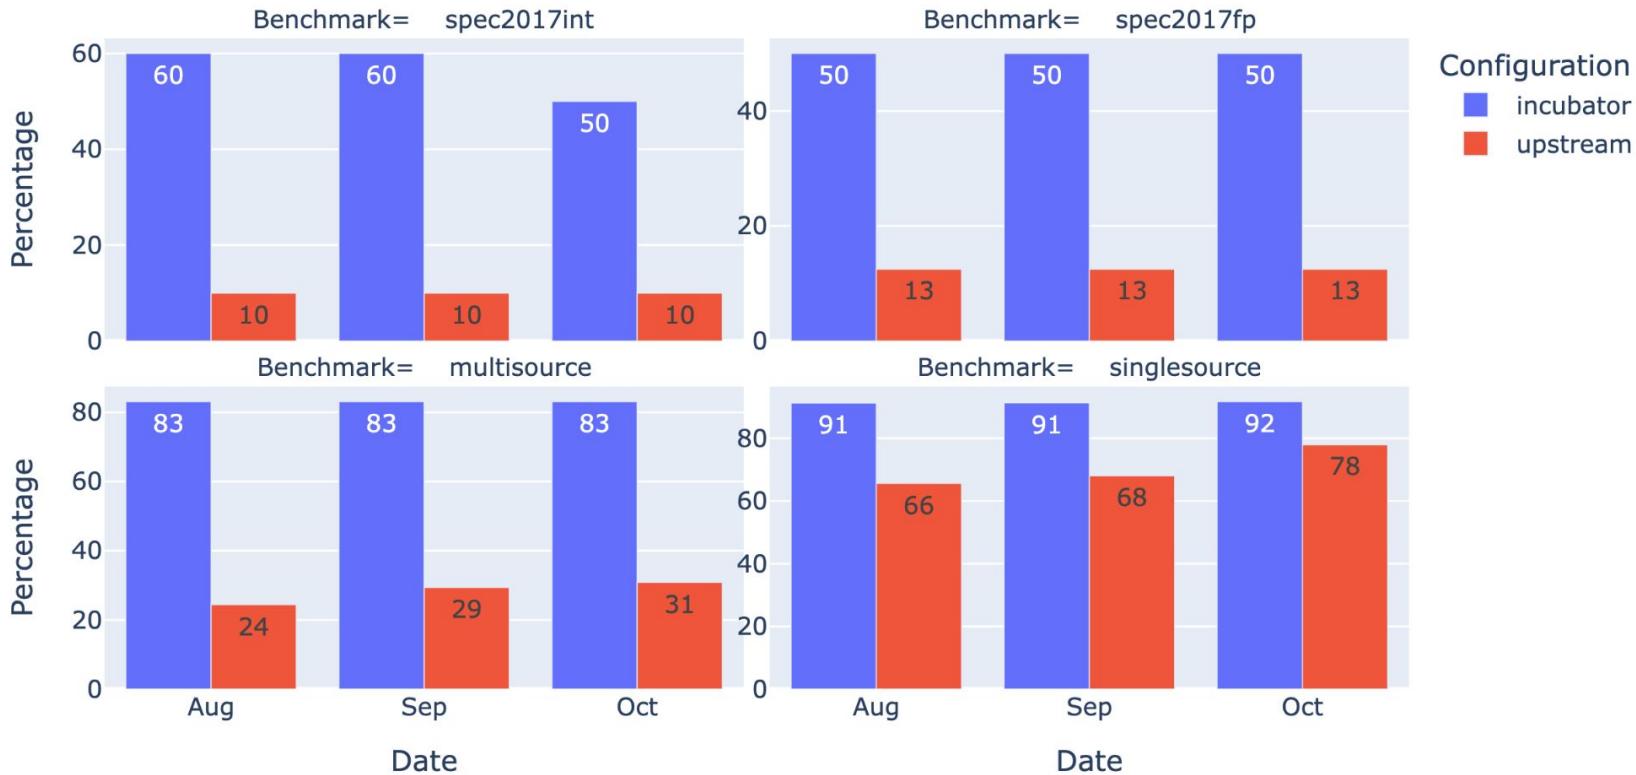
## Upstream: Grouped Commits per Month



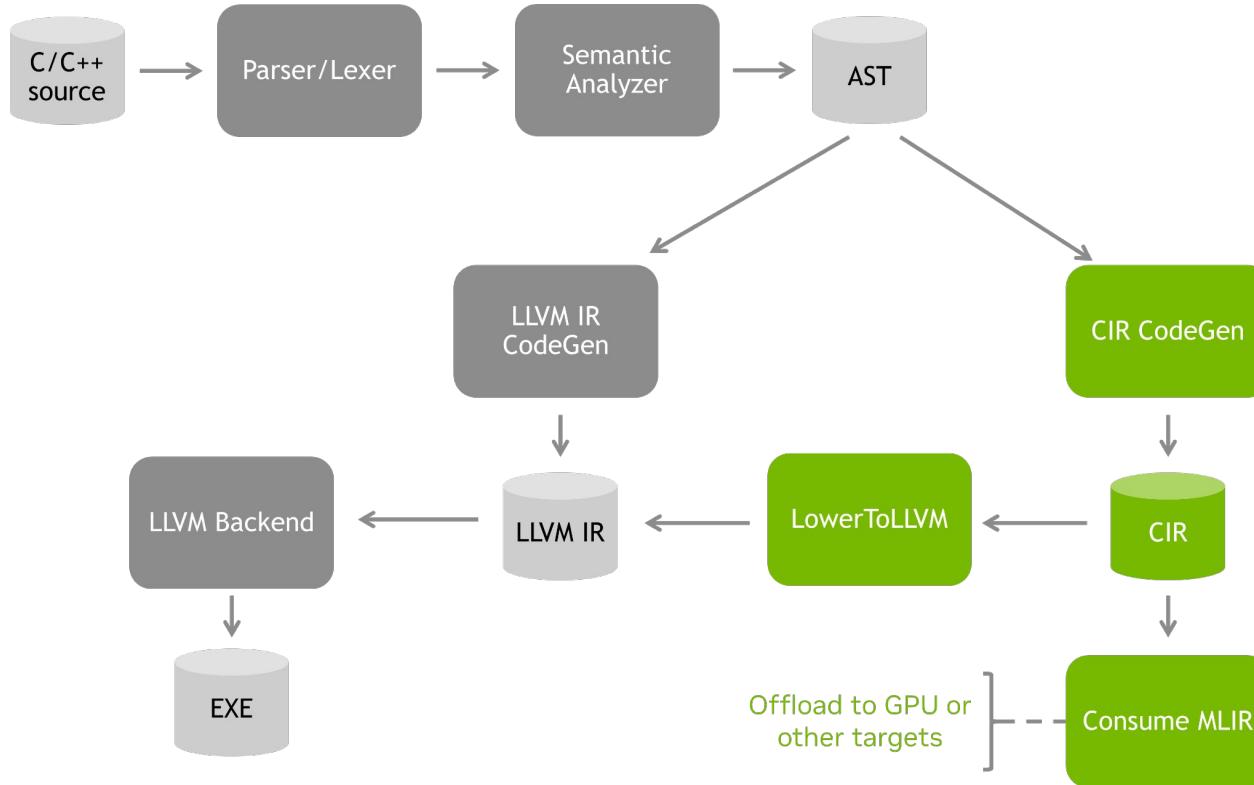
# Upstreaming progress



## Passing Benchmarks (% of completion)



# CIR Won't Always Go To LLVM IR



# Growing pains

- Build problems
  - The ClangIR build sometimes gets broken by changes in MLIR or other parts of clang
  - Sometimes the incubator code is using AST interfaces that no longer exist upstream
    - This is part of the pain of rebasing the incubator
    - It also adds some overhead to the upstreaming process
- Chasing two moving targets
  - The incubator tried to follow the structure of clang's LLVM IR codegen
    - ...at the time the incubator code was written
  - Development is still proceeding in the incubator
  - Upstreaming pulls from the incubator, but we also need to check for changes in LLVM IR codegen

# Trying to synchronize with classic codegen

- We want generated CIR to be semantically equivalent to the LLVM IR clang generates
- LLVM IR codegen is always changing
- Strategies
  - We've tried to use the same code structures whenever possible
    - CIRGenFunction looks very much like CodeGenFunction
  - Most of our tests check LLVM IR lowering from CIR and direct LLVM IR codegen
    - This will alert us when something changes in the non-CIR path
  - Eventually we'd like to run all existing codegen tests with CIR enabled

# Embarrassingly Parallel Code?

clang/lib/CIR/CodeGen/CIRGenExprScalar.cpp

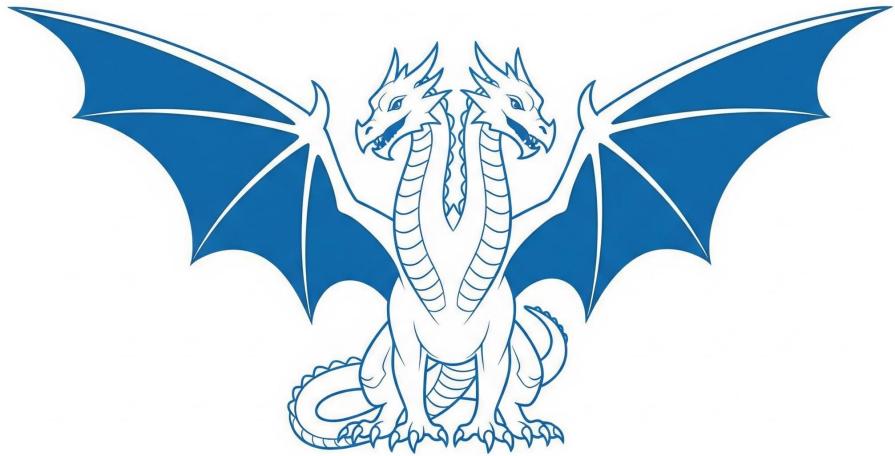
```
// Unary Operators.  
mlir::Value VisitUnaryPostDec (const UnaryOperator *e) {  
    LValue lv = cfg.emitLValue (e->getSubExpr());  
    return emitScalarPrePostIncDec (e, lv,  
cir::UnaryOpKind::Dec,  
                                    false);  
}  
mlir::Value VisitUnaryPostInc (const UnaryOperator *e) {  
    LValue lv = cfg.emitLValue (e->getSubExpr());  
    return emitScalarPrePostIncDec (e, lv,  
cir::UnaryOpKind::Inc,  
                                    false);  
}  
mlir::Value VisitUnaryPreDec (const UnaryOperator *e) {  
    LValue lv = cfg.emitLValue (e->getSubExpr());  
    return emitScalarPrePostIncDec (e, lv,  
cir::UnaryOpKind::Dec,  
                                    true);  
}  
mlir::Value VisitUnaryPreInc (const UnaryOperator *e) {  
    LValue lv = cfg.emitLValue (e->getSubExpr());  
    return emitScalarPrePostIncDec (e, lv,  
cir::UnaryOpKind::Inc,  
                                    true);  
}
```

clang/lib/CodeGen/CGExprScalar.cpp

```
// Unary Operators.  
Value *VisitUnaryPostDec (const UnaryOperator *E) {  
    LValue LV = EmitLValue (E->getSubExpr());  
    return EmitScalarPrePostIncDec (E, LV, false, false);  
}  
Value *VisitUnaryPostInc (const UnaryOperator *E) {  
    LValue LV = EmitLValue (E->getSubExpr());  
    return EmitScalarPrePostIncDec (E, LV, true, false);  
}  
Value *VisitUnaryPreDec (const UnaryOperator *E) {  
    LValue LV = EmitLValue (E->getSubExpr());  
    return EmitScalarPrePostIncDec (E, LV, false, true);  
}  
Value *VisitUnaryPreInc (const UnaryOperator *E) {  
    LValue LV = EmitLValue (E->getSubExpr());  
    return EmitScalarPrePostIncDec (E, LV, true, true);  
}
```

# Code sharing

- How do we keep CIR in sync with clang's LLVM IR codegen?
- Move common code into shared components
- Can we use templates or concepts to abstract IR-specific interfaces?



# Getting help



- Upstreaming work requires discipline
- It doesn't require prior in depth understanding
- A great way to learn!
- We've gotten a lot of help from great new contributors
- If you'd like to get involved, let us know!
- Come to the ClangIR roundtable tomorrow at 2:15



## Help Us Shape the Future

Meet us Tuesday during the LLVM Developers' Meeting 2025:

**Tuesday:** 3:00 Meet & Greet the LLVM Foundation Sponsors/California Ballroom

6:00 Evening Reception/Guildhouse



Scan the QR code to connect with us.

NVIDIA is the engine of the world's AI infrastructure. We are the world leader in accelerated computing.

Check out our open career opportunities here:

<https://www.nvidia.com/en-us/about-nvidia/careers/>



compilerjobs@meta.com