

Introducing ClangIR

High-Level IR for the C/C++ Family of Languages

Bruno Cardoso Lopes



Background


Compilation pipeline

- Multiple representations from source to machine code
- Each translation level requires **specific** information

Compilation pipeline

Progressive lowering

- Lowering: loss of information, each level is better at something



Compilation pipeline

Premature lowering

- May preclude language specific analysis & optimizations
- Reconstruction can be hard, expensive and brittle



Clang

Compiler C/C++ family of languages

- C++ is hard: more opt and analysis require richer IR
- Pipeline: C++ → AST → LLVM IR → [...] → assembly
 - AST too high level
 - LLVM IR too low level (e.g. opaque ptrs)

Clang

Why we need a new IR?

- Enable more static analysis and unlock optimization opportunities
- Success stories of high-level IRs
- Flang, Mojo, Rust, Swift, Open64's WHIRL

Clang

Reconstruction is hard


```
void f(std::vector<int> &v) {  
    v.push_back(3);  
}
```

Clang

Reconstruction is hard

-emit-llvm -O1

```
void f(std::vector<int> &v) {  
    v.push_back(3);  
}
```



<https://godbolt.org/z/zd15hK9cb>

```
define dso_local void  
%2 = getelementptr i8, ptr %1, i32 0  
%3 = load i32, ptr %2  
%4 = getelementptr i8, ptr %1, i32 1  
%5 = load i32, ptr %4  
%6 = icmp eq i32 %3, %5  
br i1 %6, label %7  
  
7:  
    store i32 3, ptr %3, %8  
    %8 = getelementptr i8, ptr %1, i32 1  
    store i32 3, ptr %8, %9  
    br label %10  
  
9:  
    %10 = load i32, ptr %1, align 8  
    %11 = ptrtoint i32 %10 to i64  
    %12 = ptrtoint i64 %11 to i64  
    %13 = sub i64 %11, %12  
    %14 = icmp eq i64 %13, 9223372036854775804  
    br i1 %14, label %15, label %16  
  
15:  
    tail call void @_ZSt20__throw_length_errorPKc(ptr @.str)  
unreachable  
  
16:  
    %17 = ashr exact i64 %13, 2  
    %18 = tail call i64 @llvm.umax.i64(i64 %17, i64 1)  
    %19 = add i64 %18, %17  
    %20 = icmp ult i64 %19, %17  
    %21 = tail call i64 @llvm.umin.i64(i64 %19,  
        i64 2305843009213693951)  
    %22 = select i1 %20, i64 2305843009213693951, i64 %21  
    %23 = icmp ne i64 %22, 0  
    tail call void @llvm.assume(i1 %23)  
    %24 = shl nuw nsw i64 %22, 2  
    %25 = tail call noalias ptr @_Znwm(i64 %24) #8  
    %26 = getelementptr inbounds i8, ptr %25, i64 %13  
    store i32 3, ptr %26, align 4  
    %27 = icmp sgt i64 %13, 0  
    br i1 %27, label %28, label %29  
  
28:  
    tail call void @_Znwm(i64 %13, i1 false)  
    br label %29  
  
29:  
    %30 = icmp eq i64 %10, null  
    br i1 %30, label %31, label %34  
    br label %32  
  
31:  
    %32 = ptrtoint ptr %5 to i64  
    %33 = sub i64 %32, %12  
    tail call void @_ZdlPvm(ptr %10, i64 %33) #9  
    br label %34  
  
34:  
    %35 = getelementptr inbounds i8, ptr %26, i64 4  
    store i32 3, ptr %35, align 8  
    store i32 3, ptr %35, align 8  
    %36 = getelementptr inbounds i32, ptr %25, i64 %22  
    store i32 3, ptr %36, align 8  
    br label %37  
  
37:  
    ret void  
}
```

ClangIR

ClangIR (CIR)

High-level IR for Clang

- Represents C/C++ closely
- Translated out of Clang's AST
- Move Clang onto the MLIR substrate
 - Use MLIR from C, C++ and extensions


ClangIR (CIR)

Open Source

- llvm-project incubator, currently being upstreamed
 - June 2022: Introductory RFC to LLVM project
 - Feb 2024: Upstream RFC in Feb 2024 (accepted)
- Github, 46 unique contributors since 2021
- Industry commitment

ClangIR (CIR)

Pipeline purview



CIR example

High-level IR for Clang

```
2
3  class A { int a; };
4  class B {
5    int b;
6    public: A *getA();
7  };
8
9  class X : public A, public B { int x; };
10 A *B::getA() { return static_cast<X*>(this); }
11
```

CIR example

High-level IR for Clang

```
2
3  class A { int a; };
4  class B {
5    int b;
6    public: A *getA();
7  };
8
9  class X : public A, public B { int x; };
10 A *B::getA() { return static_cast<X*>(this); }
11
```

Types, ABI information

```
!ty_A = !cir.struct<class "A" {!s32i}>
!ty_B = !cir.struct<class "B" {!s32i}>
!ty_X = !cir.struct<class "X" {!ty_A, !ty_B, !s32i}>
module @"sc24.cpp" attributes {
  cir.lang = #cir.lang<cxx>,
  cir.triple = "aarch64-none-linux-android24",
  ...
```

<https://godbolt.org/z/MTaPP7xdc>

CIR example

High-level IR for Clang

```
2
3  class A { int a; };
4  class B {
5    int b;
6    public: A *getA();
7  };
8
9  class X : public A, public B { int x; };
10 A *B::getA() { return static_cast<X*>(this); }
11
```

<https://godbolt.org/z/MTaPP7xdc>

Types, ABI information

```
!ty_A = !cir.struct<class "A" {!s32i}>
!ty_B = !cir.struct<class "B" {!s32i}>
!ty_X = !cir.struct<class "X" {!ty_A, !ty_B, !s32i}>
module @"sc24.cpp" attributes {
  cir.lang = #cir.lang<cxx>,
  cir.triple = "aarch64-none-linux-android24",
  ...
```

C++ idioms

```
cir.func @_ZN1B6getAsAEv(%this_param: !cir.ptr<!ty_B>) -> !cir.ptr<!ty_A> {
  %this = cir.alloca !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  cir.store %this_param, %this : !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>

}
```

CIR example

High-level IR for Clang

```
2
3  class A { int a; };
4  class B {
5    int b;
6    public: A *getA();
7  };
8
9  class X : public A, public B { int x; };
10 A *B::getA() { return static_cast<X*>(this); }
11
```

<https://godbolt.org/z/MTaPP7xdc>

Types, ABI information

```
!ty_A = !cir.struct<class "A" {!s32i}>
!ty_B = !cir.struct<class "B" {!s32i}>
!ty_X = !cir.struct<class "X" {!ty_A, !ty_B, !s32i}>
module @"sc24.cpp" attributes {
  cir.lang = #cir.lang<cxx>,
  cir.triple = "aarch64-none-linux-android24",
  ...
```

C++ idioms

```
cir.func @_ZN1B6getAsAEv(%this_param: !cir.ptr<!ty_B>) -> !cir.ptr<!ty_A> {
  %this = cir.alloca !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  cir.store %this_param, %this : !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  %b_ptr = cir.load %this : !cir.ptr<!cir.ptr<!ty_B>>, !cir.ptr<!ty_B>

}
```

CIR example

High-level IR for Clang

```
2
3  class A { int a; };
4  class B {
5    int b;
6    public: A *getA();
7  };
8
9  class X : public A, public B { int x; };
10 A *B::getA() { return static_cast<X*>(this); }
11
```

<https://godbolt.org/z/MTaPP7xdc>

Types, ABI information

```
!ty_A = !cir.struct<class "A" {!s32i}>
!ty_B = !cir.struct<class "B" {!s32i}>
!ty_X = !cir.struct<class "X" {!ty_A, !ty_B, !s32i}>
module @"sc24.cpp" attributes {
  cir.lang = #cir.lang<cxx>,
  cir.triple = "aarch64-none-linux-android24",
  ...
```

C++ idioms

```
cir.func @_ZN1B6getAsAEv(%this_param: !cir.ptr<!ty_B>) -> !cir.ptr<!ty_A> {
  %this = cir.alloca !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  cir.store %this_param, %this : !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  %b_ptr = cir.load %this : !cir.ptr<!cir.ptr<!ty_B>>, !cir.ptr<!ty_B>
  %x_ptr = cir.derived_class_addr(%b_ptr : !cir.ptr<!ty_B> nonnull) [4] -> !cir.ptr<!ty_X>
}
```

CIR example

High-level IR for Clang

```
2
3  class A { int a; };
4  class B {
5    int b;
6    public: A *getA();
7  };
8
9  class X : public A, public B { int x; };
10 A *B::getA() { return static_cast<X*>(this); }
11
```

<https://godbolt.org/z/MTaPP7xdc>

Types, ABI information

```
!ty_A = !cir.struct<class "A" {!s32i}>
!ty_B = !cir.struct<class "B" {!s32i}>
!ty_X = !cir.struct<class "X" {!ty_A, !ty_B, !s32i}>
module @"sc24.cpp" attributes {
  cir.lang = #cir.lang<cxx>,
  cir.triple = "aarch64-none-linux-android24",
  ...
```

C++ idioms

```
cir.func @_ZN1B6getAsAEv(%this_param: !cir.ptr<!ty_B>) -> !cir.ptr<!ty_A> {
  %this = cir.alloca !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  cir.store %this_param, %this : !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  %b_ptr = cir.load %this : !cir.ptr<!cir.ptr<!ty_B>>, !cir.ptr<!ty_B>
  %x_ptr = cir.derived_class_addr(%b_ptr : !cir.ptr<!ty_B> nonnull) [4] -> !cir.ptr<!ty_X>
  %a_ptr = cir.base_class_addr(%x_ptr : !cir.ptr<!ty_X>) [0] -> !cir.ptr<!ty_A>
  cir.return %a_ptr : !cir.ptr<!ty_A>
}
```

ClangIR progress

LLVM IR backend

- CIR to LLVM IR dialect pass
- Supports: x86_64, ARM64 and **SPIRV LLVM IR**
- Initial **OpenCL** support, toy **OpenMP** support
- Builds SPEC2017 C, 90% of Social App
- C++ under heavy development (WIP building libc++)

LLVM lowering

Different representation levels

```
cir.func @_ZN1B6getAsAEv(%this_param: !cir.ptr<!ty_B>) -> !cir.ptr<!ty_A> {          ClangIR
  %this = cir.alloca !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  cir.store %this_param, %this : !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  %b_ptr = cir.load %this : !cir.ptr<!cir.ptr<!ty_B>>, !cir.ptr<!ty_B>
  %x_ptr = cir.derived_class_addr(%b_ptr : !cir.ptr<!ty_B> nonnull) [4] -> !cir.ptr<!ty_X>
  %a_ptr = cir.base_class_addr(%x_ptr : !cir.ptr<!ty_X>) [0] -> !cir.ptr<!ty_A>
  cir.return %a_ptr : !cir.ptr<!ty_A>
}
```

```
define ... ptr @_ZN1B6getAsAEv(ptr %this) {          LLVM IR
entry:
  %this.addr = alloca ptr, align 8
  store ptr %this, ptr %this.addr, align 8
  %this1 = load ptr, ptr %this.addr, align 8
  %sub.ptr = getelementptr inbounds i8, ptr %this1, i64 -4
  ret ptr %sub.ptr
}
```

LLVM lowering

Different representation levels

```
cir.func @_ZN1B6getAsAEv(%this_param: !cir.ptr<!ty_B>) -> !cir.ptr<!ty_A> {          ClangIR
  %this = cir.alloca !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  cir.store %this_param, %this : !cir.ptr<!ty_B>, !cir.ptr<!cir.ptr<!ty_B>>
  %b_ptr = cir.load %this : !cir.ptr<!cir.ptr<!ty_B>>, !cir.ptr<!ty_B>
  %x_ptr = cir.derived_class_addr(%b_ptr : !cir.ptr<!ty_B> nonnull) [4] -> !cir.ptr<!ty_X>
  %a_ptr = cir.base_class_addr(%x_ptr : !cir.ptr<!ty_X>) [0] -> !cir.ptr<!ty_A>
  cir.return %a_ptr : !cir.ptr<!ty_A>
}
```

```
define ... ptr @_ZN1B6getAsAEv(ptr %this) {          LLVM IR
entry:
  %this.addr = alloca ptr, align 8
  store ptr %this, ptr %this.addr, align 8
  %this1 = load ptr, ptr %this.addr, align 8
  %sub.ptr = getelementptr inbounds i8, ptr %this1, i64 -4
  ret ptr %sub.ptr
}
```

Other Lowering

Easy to write conversions

- Built on top of MLIR
- CIR to MLIR "standard" dialects:
 - affine, arithmetic, mermen, scf, math, etc
 - Not as advanced as LLVM lowering

Tooling and Usages

Integration with existing tools

- CIR support in Compiler Explorer
- C++ lifetime analysis
 - Handle most of C++ support constructs
 - clang-tidy & clangd integration
- PoC of cross-library optimization framework

Why the HPC community
should care?

HPC & ClangIR

- Lower C/C++ extensions to MLIR
- Mix CIR with downstream and custom dialects
- High level mapping of specific C/C++ extension idioms
 - Domain specific optimizations, analysis, diagnostics
 - Avoid premature lowering

Case study

OpenMP in Clang

```
2
3 void openmp_parallel_for(int *arr, int array_size, int val)
4 {
5 #pragma omp parallel for
6     for (int i = 0; i < array_size; i++)
7         arr[i] += val;
8 }
9
```

Case study

OpenMP in Clang

```
2
3 void openmp_parallel_for(int *arr, int array_size, int val)
4 {
5 #pragma omp parallel for
6     for (int i = 0; i < array_size; i++)
7         arr[i] += val;
8 }
9
```

- Read-only variables above
- What kind of code generation we get?

Case study

OpenMP in Clang

- Read-only?
- What kind of parallelism?

```
2
3
4 define void @openmp_parallel_for(ptr %0, i32 %1, i32 %2) {
5     %arr = alloca ptr, align 8
6     %array_size = alloca i32, align 4
7     %val = alloca i32, align 4
8     store ptr %0, ptr %arr, align 8
9     store i32 %1, ptr %array_size, align 4
10    store i32 %2, ptr %val, align 4
11    call void (ptr, i32, ptr, ...) @_kmpc_fork_call(ptrnonnull @4, i32 3,
12            ptr @openmp_parallel_for_outlined,
13            ptr %array_size, ptr %arr, ptr %val)
14    ret void
15 }
```

-emit-llvm -O2 -fopenmp

Case study

OpenMP in Clang

```
define void @openmp_parallel_for(ptr %0, i32 %1, i32 %2) {
    %arr = alloca ptr, align 8
    %array_size = alloca i32, align 4
    %val = alloca i32, align 4
    store ptr %0, ptr %arr, align 8
    store i32 %1, ptr %array_size, align 4
    store i32 %2, ptr %val, align 4
    call void (ptr, i32, ptr, ...) @_kmpc_fork_call(ptrnonnull @4, i32 3,
        ptr @openmp_parallel_for_outlined,
        ptr %array_size, ptr %arr, ptr %val)
    ret void
}
```

- Unnecessary alloca's before forking

Case study

OpenMP in Clang

```
2
3 void openmp_parallel_for(int *arr, int array_size, int val)
4 {
5     #pragma omp parallel for
6         for (int i = 0; i < array_size; i++)
7             arr[i] += val;
8 }
9
```

- Function is outlined prematurely, too late for classic clang
- ClangIR: mix OpenMP + CIR
 - mem2reg remove allocas

Case study

OpenMP in Clang

```
2
3  void openmp_parallel_for(int *arr, int array_size, int val)
4  {
5      #pragma omp parallel for firstprivate(arr, array_size, val)
6      for (int i = 0; i < array_size; i++)
7          arr[i] += val;
8  }
9
```

Case study

OpenMP in Clang

```
2
3  void openmp_parallel_for(int *arr, int array_size, int val)
4  {
5      #pragma omp parallel for firstprivate(arr, array_size, val)
6      for (int i = 0; i < array_size; i++)
7          arr[i] += val;
8  }
9
```

- Work around existing compiler limitations
- No diagnostics on “writes” to those variables

Case study

OpenMP in Clang

```
2
3
4 define void @openmp_parallel_for(ptr %arr, i32 %1, i32 %2) {
5   %array_size = zext i32 %1 to i64
6   %val = zext i32 %2 to i64
7   tail call void (ptr, i32, ptr, ...) @_kmpc_fork_call(ptr @4, i32 3,
8     ptr @openmp_parallel_for_outlined,
9     i64 %array_size, ptr %arr, i64 %val)
  ret void
}
```

- Work around existing compiler limitations
- No diagnostics on “writes” to those variables

Case study

Does this happen in real code?

Case Does t

```
template<typename T>
ompBLAS_status gemm_impl(ompBLAS_handle& handle,
                          const char transa,
                          const char transb,
                          const int M,
                          const int N,
                          const int K,
                          const T& alpha,
                          const T* const A,
                          const int lda,
                          const T* const B,
                          const int ldb,
                          const T& beta,
                          T* const C,
                          const int ldc)
{
    if (M == 0 || N == 0 || K == 0)
        return 0;

    if (transa == 'T' && transb == 'N') //A(ji) * B(jk) -> C(ik)
    {
        PRAGMA_OFFLOAD("omp target teams distribute parallel for collapse(2) is_device_ptr(A, B, C)")
        for (size_t m = 0; m < M; m++)
            for (size_t n = 0; n < N; n++)
                {
```

qmcpack/src/Platforms/OMPTarget/ompBLAS.cpp

Takeaway

- Premature lowering hurts
- A higher level for C, C++ and extensions brings a clear benefit to the Clang compiler community (looking at you HPC folks)
- ClangIR is under heavy development, joins us!



Resources

- clangir.org
- Compiler explorer (ClangIR branch)
- C/C++ MLIR WG meeting monthly (1st Monday of the month)
- Discord: #clangir
- Github: <https://github.com/llvm/clangir>

Questions